



# Netezza Performance Server Health Check Report



## Document History

---


### Revision History

Revision	Date	Summary of Changes	Authors
1.0	05/08/2024	Document Creation	Sreekanth Rajan

---

### Approvals

Name	Title



## Contents

<b>1.</b>	<b>Introduction</b>	<b>4</b>
<b>2.</b>	<b>System Overview</b>	<b>4</b>
<b>2.1</b>	<b>Hardware Configuration</b>	<b>4</b>
<b>2.2</b>	<b>Software version</b>	<b>6</b>
<b>3.</b>	<b>System State</b>	<b>7</b>
<b>4.</b>	<b>System Registry Setting</b>	<b>8</b>
<b>5.</b>	<b>Performance Metrics</b>	<b>9</b>
<b>5.1</b>	<b>CPU &amp; Memory Utilization</b>	<b>9</b>
<b>5.2</b>	<b>Disk Performance Analysis</b>	<b>10</b>
<b>5.3</b>	<b>Network Usage Analysis</b>	<b>12</b>
<b>5.4</b>	<b>Database Health Check for Netezza Systems</b>	<b>14</b>
<b>6.</b>	<b>Netezza Backup Strategy</b>	<b>15</b>
<b>7.</b>	<b>Optimize Query Performance</b>	<b>16</b>



# 1. Introduction

This document's main purpose is to provide detailed steps and commands to conduct a health check on a Netezza Performance Server for the customer. Regular health checks help to ensure the system's optimal performance and proactively identify potential issues.

# 2. System Overview

In this section of System Overview, we will gather details about hardware configuration and the current Netezza software version and ensure that the system is UpToDate with patches and updates.

## 2.1 Hardware Configuration

We need to extract the details about hardware specifications in the Netezza environment and collect all the hardware-related errors and alerts from pg.log to identify any potential issues in the system.

```
[nz@c54147v1 ~]$ nzhw
Description HW ID Location  Role  State  Security
-----
Rack      1001 rack1    Active Ok   N/A
SPA       1002 spa1     Active Ok   N/A
SPU       1003 spa1.spu1 Active Online N/A
Disk      1004 spa1.disk1 Active Ok   N/A
```

```
[nz@c54147v1 ~]$ nzhw -detail
Description HW ID Location  Role  State  Security Serial number Product
Serial Version Detail
-----
-----
----
Rack      1001 rack1    Active Ok   N/A
SPA       1002 spa1     Active Ok   N/A
SPU       1003 spa1.spu1 Active Online N/A      vspu-1-1
10.0      8 CPU Cores; 15.65GB Memory; Ip Addr: 127.0.0.1; Designated Spu
Disk      1004 spa1.disk1 Active Ok   N/A      disk-1
BC17     931.51 GiB; Model ST31000640SS ; Non SED;
```

We can use the above command to display information about virtual hardware used by NPS including SPUs virtual machines and virtual disks.



```
[nz@c54147v1 ~]$ nzstats
```

Field Name	Value
Name	c54147v1.fyre.ibm.com
Description	<sys description>
Contact	<contact name>
Location	<sys location>
IP Addr	10.21.33.132
Up Time	5572 secs
Up Time Text	1 hr, 32 mins, 52 secs
Date	05-Aug-24, 01:48:06 PDT
State	8
State Text	Online
Model	
Serial Num	<serial #>
Num SFIs	0
Num SPAs	1
Num SPUs	1
Num Data Slices	1
Num Hardware Issues	0
Num Dataslice Issues	1

The PostgreSQL pg.log file is a log file that contains various messages about the server's operations. These messages include information about connections, queries, errors, and other important events. We can identify potential issues using the pg.log

```
[nz@c54147v1 postgres]$ cat /nz/kit/log/postgres/pg.log | grep -i 'error\|warning'
```

```
2024-08-05 00:17:38.505661 PDT [96184] ERROR: Attribute 'ELAPSED' not found
2024-08-05 00:19:02.394689 PDT [96231] ERROR: relation does not exist
SYSTEM.ADMIN.PG_DATABASE
2024-08-05 00:19:18.616975 PDT [96240] ERROR: relation does not exist
SYSTEM.ADMIN.PG_TABLES
2024-08-05 00:19:27.410842 PDT [96246] ERROR: Attribute 'ELAPSED' not found
2024-08-05 00:20:57.699786 PDT [96297] ERROR: relation does not exist
SYSTEM.ADMIN.PG_STATS
2024-08-05 00:24:36.066440 PDT [96353] ERROR: 'list
error ^ found "LIST" (at char 1) expecting a keyword
2024-08-05 00:25:51.525987 PDT [96353] ERROR: 'SELECT * FROM _v_qrystat
error ^ found "SELECT" (at char 26) expecting a keyword
2024-08-05 01:30:05.575591 PDT [98308] ERROR: relation does not exist
SYSTEM.ADMIN._V_SYS_RESOURCE
2024-08-05 01:30:34.983737 PDT [98323] ERROR: relation does not exist
SYSTEM.ADMIN._V_SYS_RESOURCE
2024-08-05 01:46:37.350450 PDT [98815] ERROR: relation does not exist
SYSTEM.ADMIN._V_USERS
```



## 2.2 Software version

Updating to the latest version of Netezza is crucial for maintaining a secure, efficient, and high-performing data warehouse environment. It ensures that you can take advantage of the latest technological advancements, security enhancements, and performance optimizations, while also benefiting from ongoing support and compatibility with modern systems and tools. The newer versions of Netezza often include optimizations that enhance query performance and overall system efficiency. The updates typically include patches for known vulnerabilities, reducing the risk of security breaches. The new versions address bugs and issues identified in older versions, leading to a more stable and reliable system. Before upgrading we need to ensure compatibility with the latest operating systems and other software dependencies. The latest version of Netezza Performance Server is 11.2.2.4. This version includes several new features and improvements:

- New Features:
  1. Querying data from data lakes (AWS S3).
  2. Using Kafka as a data source or data sink.
  3. Improved backup and restore capabilities, with `enablesplitdelete` enabled by default.
- Resolved Issues:
  1. Fixed issues with table broadcasts on all SPUs, host node failover, and frequent Postgres crashes.
  2. Enhanced stability and performance with critical patches.

```
[nz@c54147v1 ~]$ nzrev  
Release 11.3.0.0 [Build 4450]
```

We can use `nzrev` to check the installed Netezza build version.

```
[nz@c54147v1 ~]$ nzrev -V  
11.3.0.0-P0-F0-Bld4450
```



### 3. System State

It is very important to check system uptime and identify any recent reboots, as it is crucial to ensure that Netezza is running and available.

```
[nz@c54147v1 ~]$ uptime
02:22:17 up 2:28, 1 user, load average: 0.07, 0.02, 0.00
```

```
[nz@c54147v1 ~]$ nzstate
System state is 'Online'.
```

```
[nz@c54147v1 ~]$ nzsystem showIssues
```

Spu Partition Issues :

SPU	Partition Id	Partition Type	Status	Size (GiB)	% Used	Supporting Disks
-----	--------------	----------------	--------	------------	--------	------------------

1003 0	Data	Unknown	16	0.00	1004	
1003 100	NzLocal	Unknown	868	0.00	1004	
1003 101	Swap	Unknown	16	0.00	1004	
1003 110	Log	Unknown	1	0.00	1004	

The above command shows the issues with the system.

```
[nz@c54147v1 ~]$ nzstats
```

Field Name	Value
Name	c54147v1.fyre.ibm.com
Description	<sys description>
Contact	<contact name>
Location	<sys location>
IP Addr	10.21.33.132
Up Time	8050 secs
Up Time Text	2 hrs, 14 mins, 10 secs
Date	05-Aug-24, 02:29:25 PDT
State	8
State Text	Online
Model	
Serial Num	<serial #>
Num SFIs	0
Num SPAs	1
Num SPU's	1
Num Data Slices	1
Num Hardware Issues	0
Num Dataslice Issues	1

This command provides various statistics about the system, including the uptime.



## 4. System Registry Setting

The `nzsystem showRegistry` command in Netezza is used to display the current registry settings of the system. The registry contains various configuration settings and parameters that control the behavior and operation of the Netezza system. Running this command provides detailed information about these settings.

```
[nz@c54147v1 ~]$ nzsystem showRegistry
#
# NPS configuration registry
# Date: 05-Aug-24 02:25:41 PDT
# Revision: 11.3.0.0
```


```
startup.objCacheFiles          = 50000
startup.maxConnections         = 200
startup.enableAudit            = yes
startup.spuGlobalSharedMemoryMB = 512
```

```
[nz@c54147v1 ~]$ nzsystem showRegistry | grep -i memory
startup.spuSimMemoryMB          = 0
startup.spuGlobalSharedMemoryMB = 512
startup.spuPMSharedMemoryMB     = 0
startup.spuSharedMemoryHeaps    = 1
sysmgr.memoryErrWarning         = 1
sysmgr.memoryErrFailover        = 1
sysmgr.hostMemoryAmountWarning  = 0
sysmgr.hostMemoryAmountFailover = 0
sysmgr.hostMemoryUsageThresholdToRiseEvent = -1
sysmgr.vseriesHostMemoryUsageThresholdToRiseEvent = 99
host.snGroupMemoryMinimumPct    = -1
host.snHostMemoryQuota          = 32768
host.snHostMemoryMaxEstimate     = 3000
host.qcMaxLoadMemory            = 2000
system.bladeMemoryMB            = 400
```

Key points:

- **System:** General system settings like file paths for temporary files and databases.
- **Performance:** Settings related to system performance, such as the maximum number of concurrent queries and memory allocation.
- **Networking:** Network-related settings, including IPv6 support and client timeout configurations.
- **Security:** Security settings, such as remote connection permissions and SSL enablement.





## 5. Performance Metrics

To evaluate Netezza server performance, we should examine factors like CPU and memory usage, identify processes consuming excessive resources, and monitor disk performance using I/O statistics to detect bottlenecks. Additionally, we need to assess available storage space, and usage trends, and analyze network traffic and bandwidth utilization to identify any network issues impacting system performance.

### 5.1 CPU & Memory Utilization

We need to check current CPU and memory utilization to identify processes which are consuming more resources

```
[nz@c54147v1 ~]$ top -n 1
```

```
top - 05:01:24 up 5:07, 2 users, load average: 0.10, 0.05, 0.01
Tasks: 243 total, 1 running, 242 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.8 us, 0.8 sy, 0.0 ni, 98.5 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 15748.2 total, 6856.1 free, 860.2 used, 8031.9
buff/cache
MiB Swap: 16380.0 total, 16380.0 free, 0.0 used. 14293.4 avail Mem
```

	PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+
COMMAND											
	104656	nz	20	0	54540	4652	3720	R	12.5	0.0	0:00.02
	1	root	20	0	238564	11400	8416	S	0.0	0.1	0:03.23
systemd											
	2	root	20	0	0	0	0	S	0.0	0.0	0:00.01
kthreadd											
	3	root	0	-20	0	0	0	I	0.0	0.0	0:00.00
rcu_gp											
	4	root	0	-20	0	0	0	I	0.0	0.0	0:00.00
rcu_par_gp											
	5	root	0	-20	0	0	0	I	0.0	0.0	0:00.00
slub_flushwq											
	7	root	0	-20	0	0	0	I	0.0	0.0	0:00.00
kworker/0:0H-events_highpri											
	10	root	0	-20	0	0	0	I	0.0	0.0	0:00.00
mm_percpu_wq											
	11	root	20	0	0	0	0	S	0.0	0.0	0:00.00
rcu_tasks_rude_											

```
[nz@c54147v1 ~]$ free -m
```

	total	used	free	shared	buff/cache
available					
Mem:	15748	872	6843	264	8032
14280					
Swap:	16379	0	16379		



## 5.2 Disk Performance Analysis

In this section, we will conduct a comprehensive analysis of disk performance. The process will include the following steps:

1. Monitor Disk I/O Statistics:
  - Track key metrics such as read/write speeds, latency, and queue length.
  - Use tools like iostat, vmstat, or dstat to gather detailed statistics.
  - Identify patterns or anomalies that could indicate performance issues or inefficiencies.
2. Identify Bottlenecks:
  - Analyze the gathered data to pinpoint areas where disk performance is suboptimal.
  - Look for signs of high latency or consistently high queue lengths, which may suggest bottlenecks.
3. Evaluate Storage Space and Usage Trends:
  - Assess the current storage capacity and usage patterns.
  - Use tools like df and du to evaluate available storage and identify trends over time.
4. Plan for Capacity Expansion:
  - Based on the analysis of usage trends, determine if there is a need for additional disk space.
  - Develop a strategy to increase disk capacity as needed to ensure optimal performance and prevent future bottlenecks.

This analysis will help optimize disk performance, ensuring the system runs efficiently and is prepared for future storage demands.

```
[nz@e1n1-npshost ~]$ iostat -x 1 5
Linux 3.10.0-1160.53.1.el7.x86_64 (e1n1)      08/05/2024      _x86_64_
(32 CPU)

avg-cpu:  %user   %nice %system %iowait  %steal   %idle
           6.35    0.06   3.62    0.31    0.00   89.66

Device:            rrqm/s   wrqm/s     r/s     w/s    rkB/s    wkB/s avgrq-sz
avgqu-sz   await  r_await  w_await  svctm  %util
nvme0n1          93.58     0.00    4.68    5.93   566.54   706.38   239.78
0.00    0.17    0.22    0.13    0.50    0.53
nvme1n1          93.58     0.00    4.65    5.83   566.14   718.82   245.33
0.00    0.17    0.22    0.13    0.50    0.53
nvme3n1          93.58     0.00    4.65    7.34   566.40   757.02   220.68
0.00    0.19    0.30    0.12    0.55    0.66
nvme2n1          93.58     0.00    4.65    6.00   566.29   721.81   241.83
0.00    0.17    0.23    0.13    0.52    0.56
sda              0.42    17.67    1.08  123.92    39.44   695.29    11.76
0.10    0.84    0.77    0.84    0.03    0.39
```

```
md0          0.00      0.00      0.00      0.00      0.00      0.00      0.00      31.51
0.00      0.00      0.00      0.00      0.00      0.00
dm-0         0.00      0.00      0.00      0.00      0.00      0.00      0.00      35.74
0.00      0.15      0.14      2.05      0.07      0.00
```

```
avg-cpu:  %user   %nice %system %iowait  %steal   %idle
           2.94    0.00    1.63    0.09    0.00   95.34
```

```
Device:          rrqm/s   wrqm/s     r/s     w/s    rkB/s    kB/s avgrq-sz
avgqu-sz   await r_await w_await  svctm  %util
nvme0n1      0.00     0.00    20.00    0.00   980.00    0.00    98.00
0.00      0.15    0.15    0.00   0.25   0.50
nvme1n1      0.00     0.00    20.00    1.00   980.00   252.00   117.33
0.00      0.14    0.15    0.00   0.24   0.50
nvme3n1      0.00     0.00    20.00    3.00   980.00    12.00    86.26
0.00      0.09    0.10    0.00   0.35   0.80
nvme2n1      0.00     0.00    20.00    2.00   980.00   256.00   112.36
0.00      0.18    0.20    0.00   0.32   0.70
sda          0.00     0.00    38.00    6.00  1773.00    28.00    81.86
0.01      0.32    0.37    0.00   0.32   1.40
md0          0.00     0.00    0.00    0.00    0.00    0.00    0.00
0.00      0.00    0.00    0.00    0.00    0.00
dm-0         0.00     0.00    0.00    0.00    0.00    0.00    0.00
0.00      0.00    0.00    0.00    0.00    0.00
```

```
avg-cpu:  %user   %nice %system %iowait  %steal   %idle
           5.04    0.00    2.07    0.00    0.00   92.90
```

```
Device:          rrqm/s   wrqm/s     r/s     w/s    rkB/s    kB/s avgrq-sz
avgqu-sz   await r_await w_await  svctm  %util
nvme0n1      0.00     0.00    0.00    5.00    0.00    92.00    36.80
0.00      0.00    0.00    0.00   0.80   0.40
nvme1n1      0.00     0.00    0.00    6.00    0.00   272.00    90.67
0.00      0.17    0.00    0.17   0.50   0.30
nvme3n1      0.00     0.00    0.00   13.00    0.00   392.00    60.31
0.00      0.00    0.00    0.00   0.38   0.50
nvme2n1      0.00     0.00    0.00    1.00    0.00    4.00    8.00
0.00      0.00    0.00    0.00   1.00   0.10
sda          0.00    76.00    0.00  750.00    0.00  3536.00    9.43
1.77      2.35    0.00    2.35   0.02   1.80
md0          0.00     0.00    0.00    0.00    0.00    0.00    0.00
0.00      0.00    0.00    0.00    0.00    0.00
dm-0         0.00     0.00    0.00    0.00    0.00    0.00    0.00
0.00      0.00    0.00    0.00    0.00    0.00
```

```
[nz@e1n1-npshost ~]$ df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
overlay	98G	6.7G	91G	7%	/
tmpfs	95G	0	95G	0%	/dev
ips	5.9T	3.1T	2.9T	53%	/nz
/dev/sda2	98G	6.7G	91G	7%	
/etc/hosts					
/dev/sda9	195G	108G	77G	59%	
/etc/resolv.conf.upstream					
tmpfs	95G	3.7M	95G	1%	
/host/run					
tmpfs	19G	0	19G	0%	
/host/run/user/0					
shm	152G	551M	152G	1%	
/dev/shm					
/dev/sda5	38G	20G	17G	55%	
/var/lib/sedssupport					
tmpfs	64M	13M	52M	19%	/run
tmpfs	64M	0	64M	0%	
/run/lock					
tmpfs	64M	0	64M	0%	
/var/log/journal					
tmpfs	71G	2.9G	68G	5%	/tmp
10.231.121.246:/nfs_kbbnzbackup06_backup	79T	38T	42T	48%	
/mnt/ipsbackups					
10.231.121.246:/nas_kb_mx_sd_bo_edw_archive	80T	68T	13T	85%	
/mnt/edw_archive					

## 5.3 Network Usage Analysis

In this section, we will perform a detailed analysis of network usage to ensure optimal performance. The process will include the following steps:

- Analyze Network Traffic:**
  - Monitor network traffic patterns using tools such as ifstat, tcpdump, or ntop.
  - Collect data on packet flow, connection types, and traffic sources/destinations.
  - Identify any unusual or excessive traffic that may indicate network congestion or security issues.
- Evaluate Bandwidth Utilization:**
  - Measure current bandwidth usage to determine if the network is operating within its capacity.
  - Use tools like NetFlow, SNMP, or bmon to gain insights into bandwidth allocation and consumption.
  - Identify peak usage times and evaluate if additional bandwidth is required to accommodate demand.



### 3. Identify Network-Related Issues:

- Analyze data for signs of latency, packet loss, or jitter, which can affect network performance.
- Investigate any recurring issues, such as slow connection speeds or frequent disconnects, to determine root causes.

By thoroughly analyzing network traffic and bandwidth utilization, we can identify potential issues affecting performance and implement solutions to enhance network efficiency and reliability.

```
[nz@e1n1-npshost ~]$ ifstat -a
```

```
#kernel
```

Interface	RX Pkts/Rate	TX Pkts/Rate	RX Data/Rate	TX
Data/Rate				
Coll/Rate	RX Errs/Drop	TX Errs/Drop	RX Over/Rate	TX
lo	2855M 0	2855M 0	4242M 0	4242M 0
	0 0	0 0	0 0	0 0
mgt0	2468M 0	2212M 0	494728K 0	2506M 0
	0 0	0 0	0 0	0 0
mgt1	47466K 0	3816K 0	1960M 0	3960M 0
	0 20889	0 0	0 0	0 0
fab0	1024M 0	3794M 0	1249M 0	2346M 0
	0 2	31 0	0 0	0 0
fab1	2521M 0	703378K 0	3036M 0	923449K 0
	0 2	0 0	0 0	0 0
mgt-br0	2278M 0	2155M 0	3069M 0	2950M 0
	0 0	0 0	0 0	0 0
fbond	3546M 0	203133K 0	795332K 0	3269M 0
	0 8	31 0	0 0	0 0
fbond.4079	284174K 0	2620M 0	2397M 0	1475M 0
	0 0	0 0	0 0	0 0
fab-br0	2492M 0	202769K 0	4034M 0	1153M 0
	0 1141	0 0	0 0	0 0
docker0	39387K 0	40658K 0	3299M 0	1269M 0
	0 0	0 0	0 0	0 0
fbond.4080	344017K 0	3107M 0	2400M 0	1925M 0
	0 9593K	0 0	0 0	0 0
fbond.3121	4089M 0	2739M 0	2669M 0	1471M 0
	0 22696K	0 0	0 0	0 0
veth9aba08d	10765K 0	15606K 0	1618M 0	1820M 0
	0 0	0 0	0 0	0 0
usb0	220305 0	0 0	30784K 0	0 0
	0 0	0 0	0 0	0 0
veth13a0067	1037K 0	969879 0	105805K 0	138253K 0
	0 0	0 0	0 0	0 0

## 5.4 Database Health Check for Netezza Systems

Database health checks are crucial for maintaining the performance and efficiency of Netezza systems. A thorough health check involves analysing various aspects of the database to ensure optimal performance and prevent potential issues. Below are key considerations and steps to perform a comprehensive database health check:

### 1. Analyze Database Size and Growth Patterns

- **Monitor Database Size:** Regularly check the overall database size to understand its current state and anticipate future storage needs.
- **Growth Trends:** Analyze historical data to identify growth patterns. This helps in forecasting storage requirements and planning for capacity upgrades.
- **Capacity Planning:** Ensure the system has sufficient resources to accommodate anticipated growth without performance degradation.

### 2. Identify Large Tables and Indexes

- **Large Tables:** Identify tables that are disproportionately large and could impact performance. These tables might require special attention or optimization.
- **Index Management:** Locate large or unused indexes that might be consuming unnecessary resources. Review and optimize index usage to improve query performance.

### 3. Analyze Query Execution Time

- **Monitor Query Performance:** Identify slow-running queries that could be impacting overall system performance. Regular monitoring helps in pinpointing queries that need optimization.
- **Explain Plans:** Use explain plans to gain insights into query execution paths and identify bottlenecks. This tool is essential for diagnosing performance-related issues and making informed optimization decisions.

### 4. Check Table and Index Statistics

- **Statistics Collection:** Ensure that statistics for tables and indexes are up-to-date. Accurate statistics are vital for the query optimizer to make efficient execution plans.
- **Regular Updates:** Implement regular schedules for updating statistics, especially after large data changes, to maintain query performance.

## 5. Perform Regular Maintenance Tasks

- **Grooming Tasks:** Regularly perform grooming tasks to reclaim space from deleted or outdated data. This includes tasks like table reorganization and index rebuilding.
- **Maintenance Schedule:** Establish a routine maintenance schedule to ensure all tasks are performed consistently, preventing performance degradation over time.

## 6. Netezza Backup Strategy

A comprehensive Netezza backup strategy is crucial for ensuring data integrity, availability, and disaster recovery. As part of a health check report, the backup strategy should include regular backups, monitoring, and verification processes.

The **nzbackup** command in Netezza is used to back up the database. The **-history** option is particularly useful for viewing the history of backup operations. This includes details such as the start and end times of each backup, the status of the backup, and the size of the backup files.

By using the **nzbackup -history** command, you can keep track of your backup operations and ensure that your data is consistently protected

```
[nz@elnl-npshost ~]$ nzbackup -history
```

```
DBNAME 20240724172935 6 DIFF COMPLETED 2024-07-29 19:49:50
backupsvr.149652.2024-07-29.log
DBNAME 20240731173148 1 FULL COMPLETED 2024-07-31 10:31:48
backupsvr.156836.2024-07-31.log
DBNAME 20240731173148 2 DIFF COMPLETED 2024-07-31 19:36:34
backupsvr.150034.2024-07-31.log
DBNAME 20240731173148 3 DIFF COMPLETED 2024-08-01 19:33:29
backupsvr.8330.2024-08-01.log
DBNAME 20240731173148 4 DIFF COMPLETED 2024-08-02 19:46:23
backupsvr.60052.2024-08-02.log
DBNAME 20240731173148 5 DIFF COMPLETED 2024-08-04 19:42:11
backupsvr.101244.2024-08-04.log
DBNAME 20240731173148 6 DIFF COMPLETED 2024-08-05 19:50:11
backupsvr.140180.2024-08-05.log
```

- **Completed:** Indicates that the backup was successful.
- **Failed:** Indicates that the backup was not completed successfully. Further investigation into logs or error messages would be needed to determine the cause of failure.



## 1. Backup Types

Netezza supports several types of backups:

- **Full Backups:** Capture the entire database.
- **Incremental Backups:** Capture only the changes since the last backup.
- **Differential Backups:** Capture the changes since the last full backup.

## 2. Backup Schedule

Establish a regular backup schedule:

- **Daily Backups:** Perform incremental or differential backups daily to ensure that recent changes are captured.
- **Weekly Full Backups:** Perform full backups weekly to ensure you have a complete snapshot of your database.

### 3. Backup Retention Policy

Define how long you will keep each type of backup:

- **Daily Incremental/Differential Backups:** Retain for one week.
- **Weekly Full Backups:** Retain for one month.
- **Monthly Backups:** Retain for one year or as required by business needs.
- The following backups are recommended every week:
  - **nzhostbackup** (requires a system Pause)
  - **nzbackup -globals -dir /path**

## 7. Optimize Query Performance

This section discusses optimizing Netezza query performance, which involves steps and best practices to ensure that queries run efficiently and use system resources best. Here's a comprehensive approach to analyze and optimize Netezza query performance:

- Check active queries using `_v_qrystat` to identify long-running or resource-intensive queries.

```
[nz@c54147v1 ~]$ nzsqli -c " select * from _v_grystat "
```

```

QS_SESSIONID | QS_PLANID | QS_CLIENTID | QS_CLIIPADDR | QS_SQL |
QS_STATE | QS_TSUBMIT | QS_TSTART | QS_PRIORITY | QS_PRITXT |
QS_ESTCOST | QS_ESTDISK | QS_
ESTMEM | QS_SNIPPETS | QS_CURSNIPT | QS_RESROWS | QS_RESBYTES

```

(0 rows)





- Use `_v_qryhist` to review past query performance and identify recurring issues.

```
[nz@c54147v1 ~]$ nzsqli -c " select * from _v_qryhist "
```

```
QH_SESSIONID | QH_PLANID | QH_CLIENTID | QH_CLIIPADDR | QH_DATABASE |  
QH_USER | QH_SQL  
| QH_TSUBMIT | QH_TSTART | QH_TEND |  
QH_PRIORITY | QH_PRITXT | QH_ESTCOST | QH_ESTDISK | QH_ES  
TMEM | QH_SNIPPETS | QH_SNPTSDONE | QH_RESROWS | QH_RESBYTES |  
QH_CLIENT_USER_ID | QH_CLIENT_APPLICATION_NAME |  
QH_CLIENT_WORKSTATION_NAME | QH_CLIENT_ACCOUN  
TING_STRING
```

```
-----+-----+-----+-----+-----+-----+-----  
  
62331 | 1 | 6 | 10.21.33.132 | SYSTEM | ADMIN | SELECT datasliceid,  
COUNT(*) as rows_per_slice FROM sales GROUP BY dataslice  
id ORDER BY rows_per_slice DESC | 2024-08-06 05:35:37 | 2024-08-06 05:35:37 | 2024-  
08-06 05:35:45 | 3 | normal | 0 | 1 |  
1 | 1 | 1 | 0 | 0 | | |  
62331 | 2 | 6 | 10.21.33.132 | SYSTEM | ADMIN | INSERT INTO sales  
(sale_id, customer_id, sale_date, amount) VALUES (1, 101,  
'2024-08-01', 150.75) | 2024-08-06 05:36:51 | 2024-08-06 05:36:51 | 2024-08-06  
05:36:52 | 3 | normal | 0 | 0 |
```

- Use `EXPLAIN` to review the query execution plan.

```
SYSTEM.ADMIN(ADMIN)=> EXPLAIN SELECT * FROM sales WHERE sale_date >  
'2023-01-01' ;
```

NOTICE: QUERY PLAN:

QUERY PLANTEXT:

Sequential Scan table "SALES" (cost=0.0..0.0 rows=5 width=20 conf=80)

- Ensure tables are evenly distributed across SPUs to avoid data skew.

```
SYSTEM.ADMIN(ADMIN)=> SELECT datasliceid, COUNT(*) as rows_per_slice FROM  
sales GROUP BY datasliceid ORDER BY rows_per_slice DESC ;  
DATASLICEID | ROWS_PER_SLICE
```

```
-----+-----  
1 | 5
```



## 7.1 Zone Maps for Query Performance

Zone maps in Netezza are an essential feature for query performance optimization, particularly in large-scale databases. They help to minimize disk I/O by allowing the database engine to skip over large sections of data that are irrelevant to the query, thus speeding up query execution times. Zone maps are metadata structures that store the minimum and maximum values for data stored in each extent of a table. When a query is executed, Netezza uses these zone maps to determine which extents need to be scanned and which can be skipped, based on the query predicates. Zone maps are automatically created and maintained by Netezza when data is loaded into a table

### Best practices for effective utilization of Zone map

- We need to choose the Right Distribution Key to distribute data evenly across all data slices. A poorly chosen distribution key can lead to data skew, reducing the effectiveness of zone maps.
- For very large tables, consider partitioning them by commonly filtered columns.
- If possible, avoid applying functions to columns in the WHERE clause that benefit from zone maps, as this can prevent the zone map from being used.
- Periodically groom your tables to reorganize data and optimize the effectiveness of zone maps.
- To verify all the columns on which Netezza can create zone maps, you can use the `nz_zonemap` utility from the SQL toolkit. This toolkit includes various useful scripts and commands to assist in optimizing Netezza performance, including identifying columns that are candidates for zone maps.
- Once the SQL toolkit is installed, you can use the `nz_zonemap` script to list the columns that can have zone maps created on them

```
nz_zonemap DATABASENAME TABLENAME Distributed Column name
```

## Key Points:

- Prefer hash joins for large datasets and nested loop joins for smaller ones.
- Ensure that join keys are properly indexed and distributed.
- Regularly groom tables to reclaim space from deleted rows and reorganize data.
- Choose a column that evenly distributes data across all SPUs
- Choose columns that are frequently used in large table joins to get collocated joins
- Use the same data type for the join columns
- Choose a single column – Avoid multi-column distributions if possible
- Choose RANDOM distribution for small dimensions and lookup tables only
- When unloading large data sets use EXTERNAL TABLES instead of nzsqli
- Regularly check for data skew and redistribute tables as necessary.
- Running statistics in Netezza helps optimize query performance by updating the metadata about the tables and columns, such as the number of rows, distinct values, and data distribution

```
[nz@c54147v1 ~]$ nzsqli -c " GENERATE STATISTICS ON sales; "  
GENERATE STATISTICS
```

- Regularly update statistics, especially after significant data changes (e.g., bulk inserts, updates, or deletes).